



MSc in AI
NCSR Demokritos - University of Piraeus

Course: **Machine Learning**

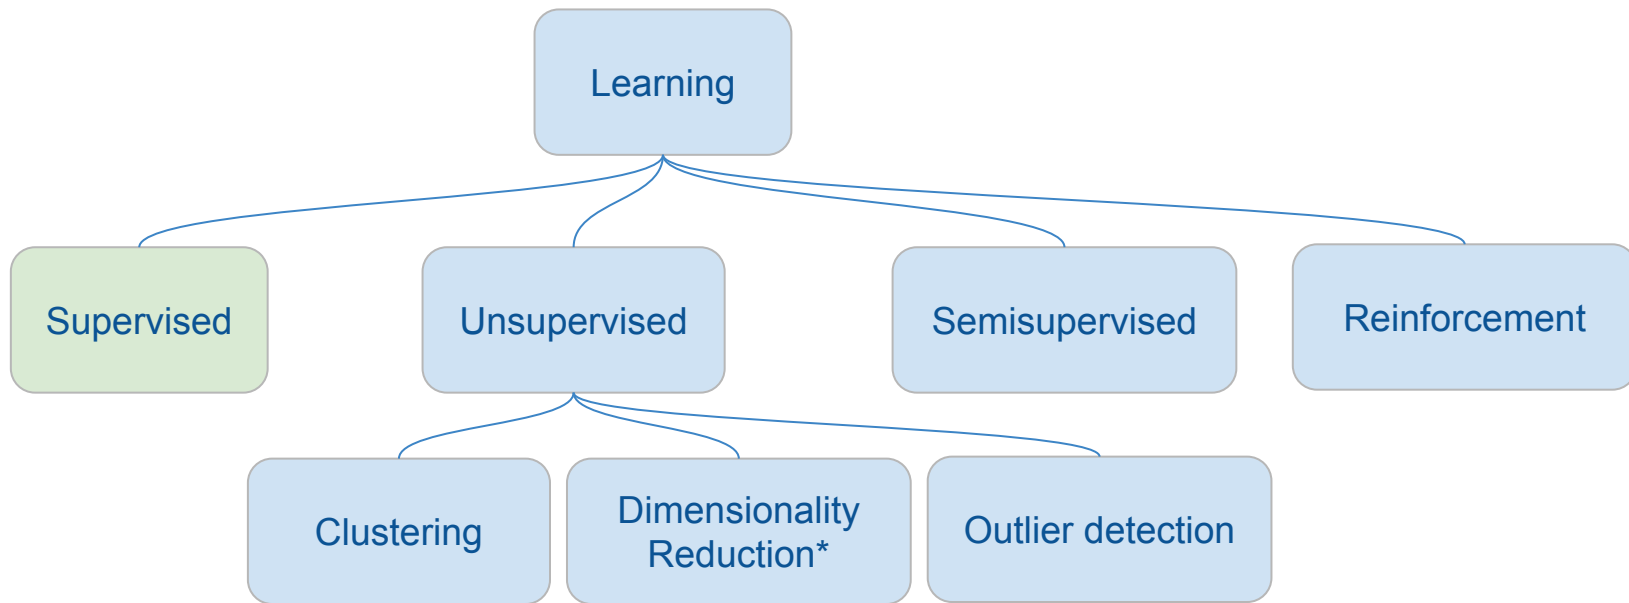
Lesson 2

Linear Regression

Theodoros Giannakopoulos



Types of learning



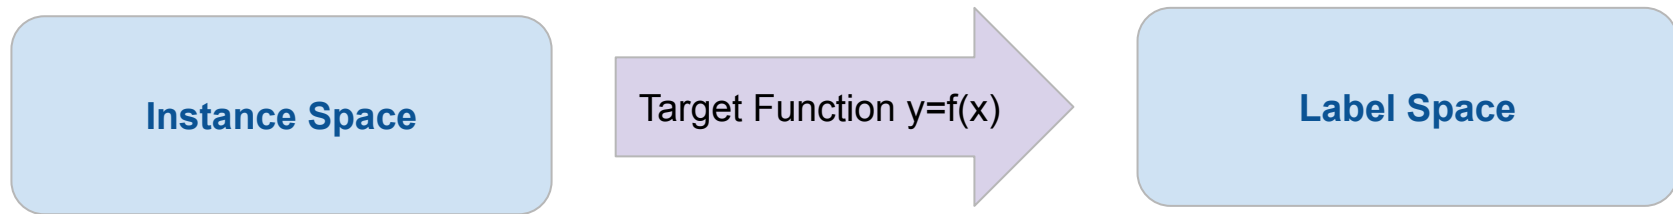
Supervised Learning

- Given data and correct output, try to find a relationship between data and output

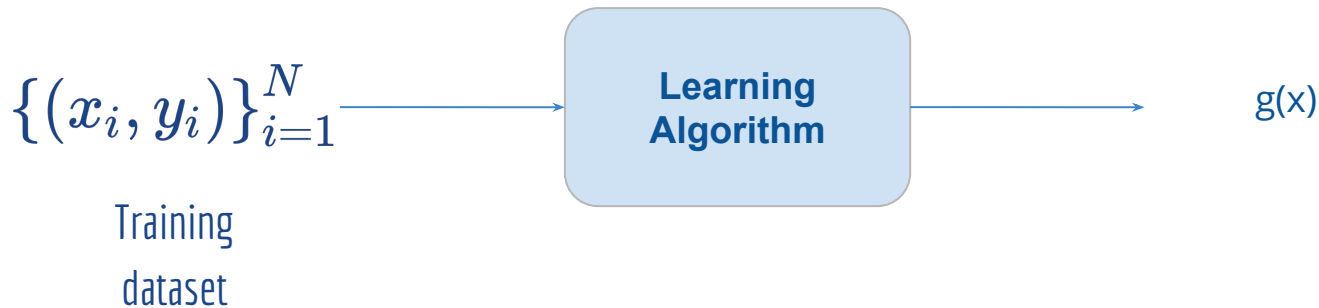
$$\{(x_i, y_i)\}_{i=1}^N$$

- N labelled examples or instances or feature vectors
- x_i feature vector $\in \mathbb{R}^D$
- D: dimensionality
- $x_i(j)$ j-th feature, $j=1, \dots, D$
- y_i label $\in \{w_1, w_2, \dots, w_c\}$ set of **classes** OR \mathbb{R} (regression)

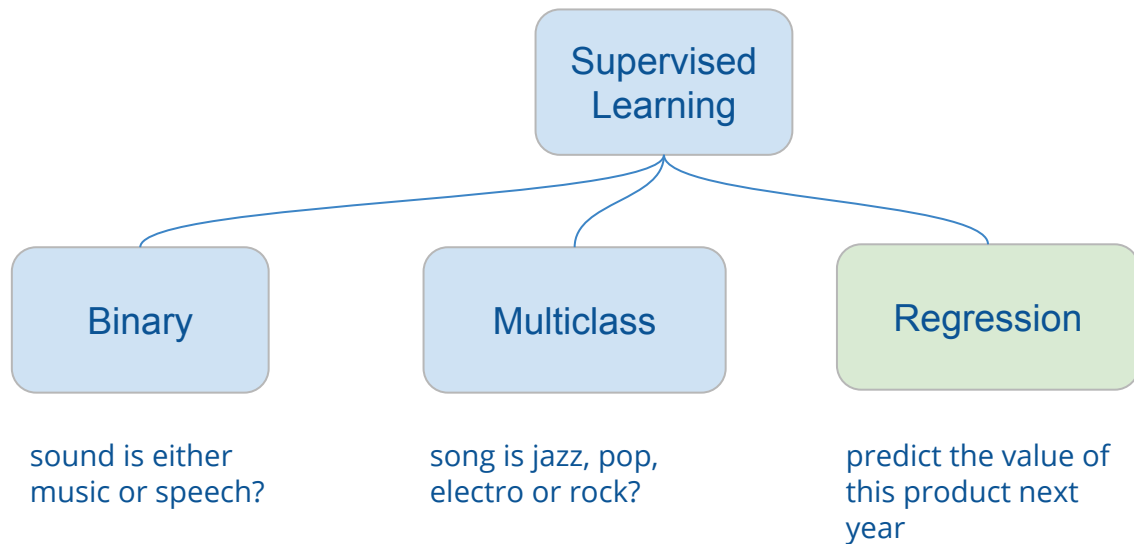
(Supervised) learning of the target function from training data



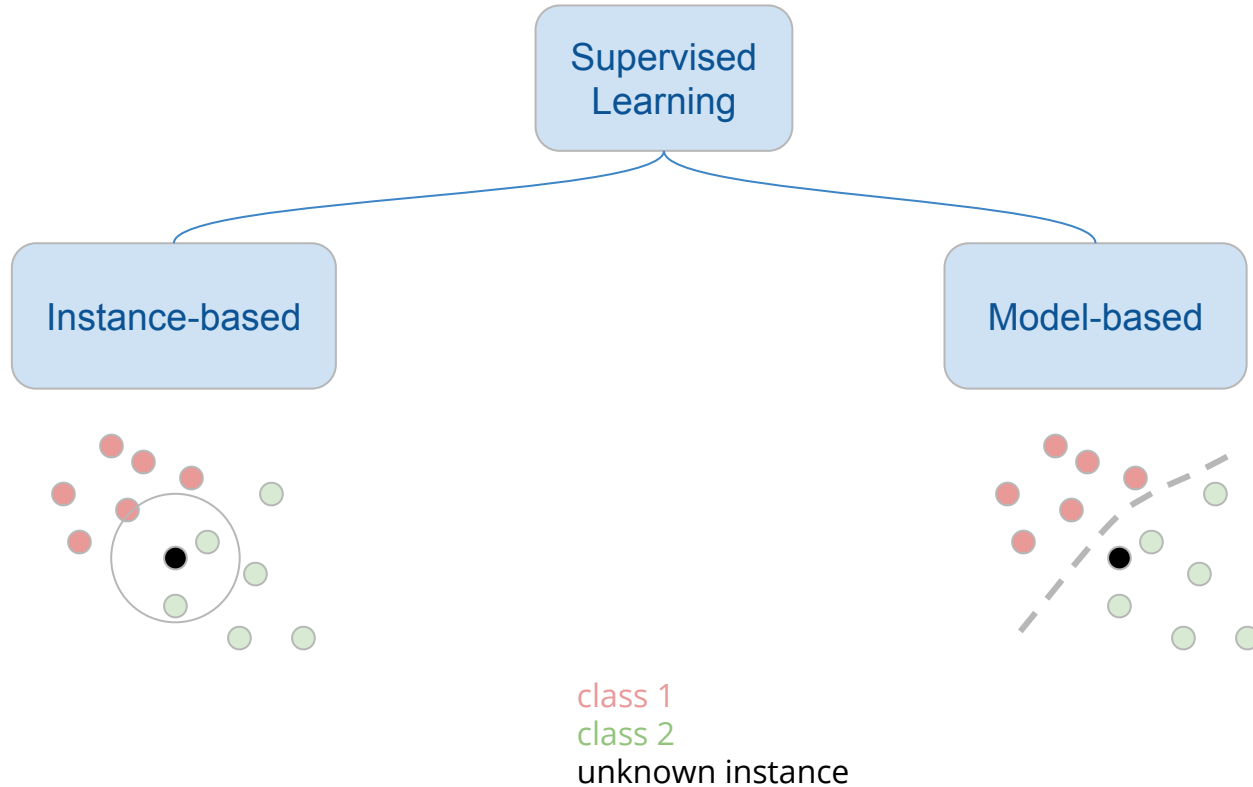
The learning algorithm does not know $f(x)$, it only “sees” some examples of it, though the training data



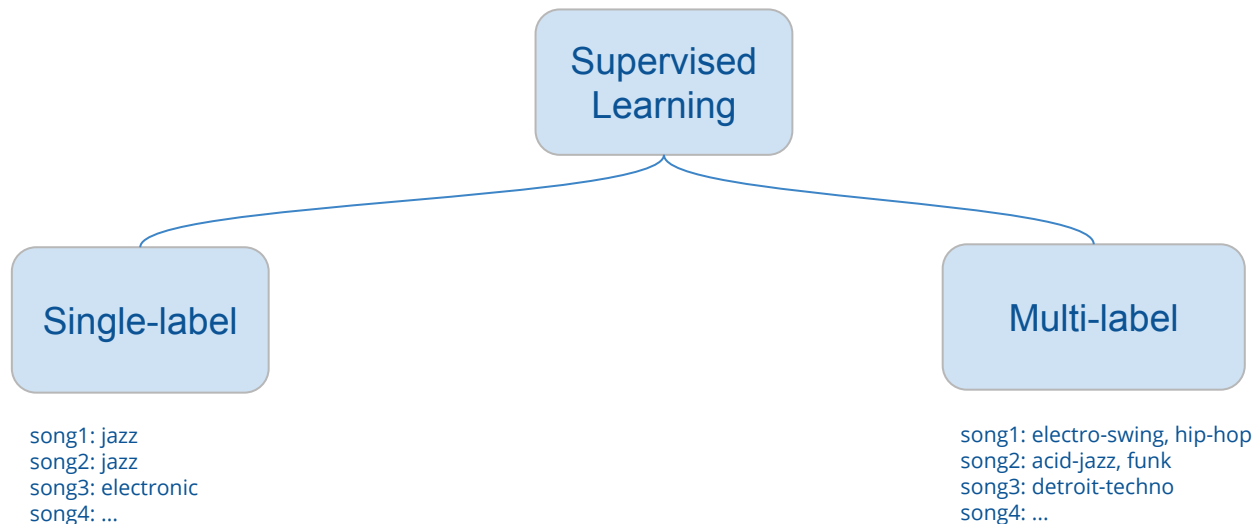
Types of supervised learning (1/3)



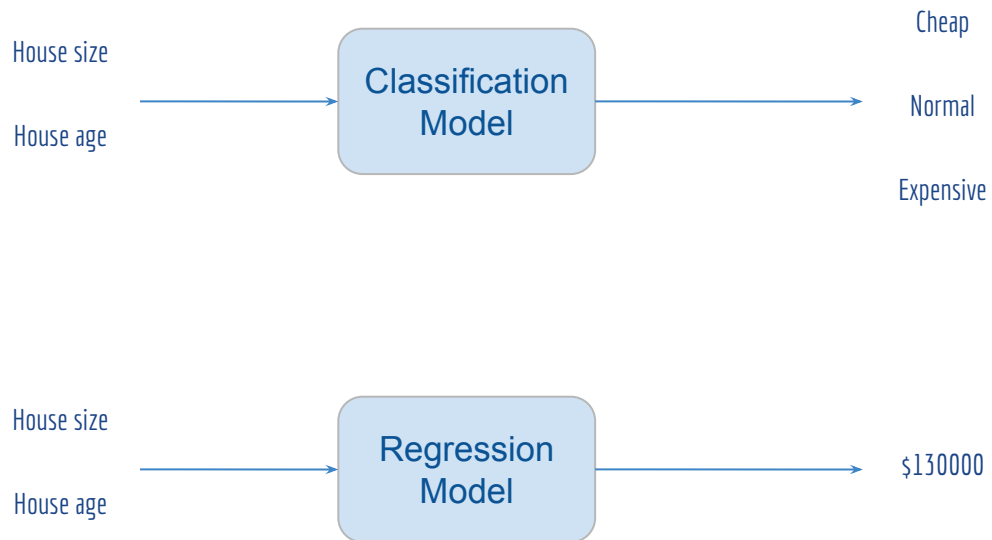
Types of supervised learning (2/3)




Types of supervised learning (3/3)




Classification Vs Regression




Linear Regression: train (1D)



➡ Training ➡




Linear Regression: test (1D)



Test




$$x = x_{\text{test}}$$



Linear Regression: problem formulation

$$f_{w,b}(\mathbf{x}) = \mathbf{w}\mathbf{x} + b$$

- \mathbf{x} : D-dimensional feature vector
- \mathbf{w} : D-dimensional vector
- b : bias
- y : target value (ground truth)



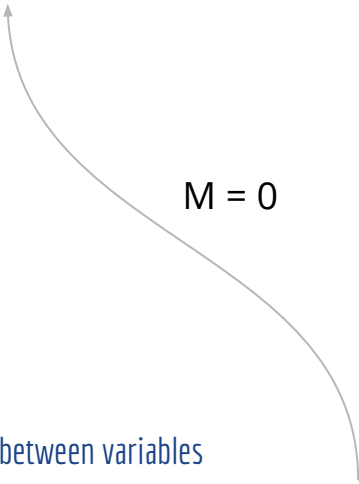
Regression: $y \in \mathbb{R}$

Linear Regression: problem formulation

$$f_{w,b}(x) = wx + b$$

- x : D-dimensional feature vector
- w : D-dimensional vector
- b : bias
- y : target value (ground truth)
- Model f is parametrized wrt w and b
- Note on linearity:
 - Linear regression can also handle **non-linear** relationships between variables
 - f can also be nonlinear to x and still be linear to the w coefficients: $w_0 + w_1x + w_2x^2 + \dots + w_Mx^M$
 - (here, we use b as w_0)
 - Don't be confused by the term "linear"!

$M = 0$



Linear Regression: problem formulation

$$f_{\mathbf{w},b}(\mathbf{x}) = \mathbf{w}\mathbf{x} + b$$

$f_{\mathbf{w},b}$ models relationship between features \mathbf{x} and target value y

$$y = f + \varepsilon$$

Goal of a regression learning algorithm: estimate **parameters** \mathbf{w} and b

(so that h makes a good prediction), i.e. such as:

$$y_i \approx f(\mathbf{x}_i)$$

Linear Regression: how to estimate w and b ?

$$f_{w,b}(\mathbf{x}) = \mathbf{w}\mathbf{x} + b$$

$$J(\mathbf{w}, b) = \frac{1}{2N} \sum_{i=0}^N (f_{w,b}(x_i) - y_i)^2$$

Cost function: measures the error between true and predicted values

Loss function: a measure of penalty for misclassification of each example i

In linear regression, cost function is the average loss (also called empirical risk)

Linear Regression: how to estimate w and b ?

$$f_{w,b}(\mathbf{x}) = \mathbf{w}\mathbf{x} + b$$

$$J(\mathbf{w}, b) = \frac{1}{2N} \sum_{i=0}^N (f_{w,b}(x_i) - y_i)^2$$

Goal:

Cost function: measures the error between true and predicted values

minimize J to find w and b

Loss function: a measure of penalty for misclassification of each example i

In linear regression, cost function is the average loss (also called empirical risk)

Linear Regression: how is J minimized?

$$J(\mathbf{w}, b) = \frac{1}{2N} \sum_{i=0}^N (f_{\mathbf{w},b}(x_i) - y_i)^2$$

Gradient Descent:

- A generic optimization algorithm, not just for LR
- Start with some w, b
- Keep changing each param w (b, w_1, w_2, \dots)

$$w := w - \alpha \frac{\partial J(w, b)}{\partial w}$$

Linear Regression: how is J minimized? (GD for 2 params)

$$J(w_1, b) = \frac{1}{2N} \sum_{i=0}^N (f_{w_1, b}(x_i) - y_i)^2 \quad f_{w_1, b} = w_1 x + b$$

Gradient Descent:

- Select a random value for w_1 and b
- Until convergence (or for a max number of epochs):

$$w_1 := w_1 - \alpha \frac{\partial J(w_1, b)}{\partial w_1}$$

$$b := b - \alpha \frac{\partial J(w_1, b)}{\partial b}$$

Linear Regression: how is J minimized? (GD for 2 params)

$$J(w_1, b) = \frac{1}{2N} \sum_{i=0}^N (f_{w_1, b}(x_i) - y_i)^2 \quad f_{w_1, b} = w_1 x + b$$

Gradient Descent:

- Select a random value for w_1 and b
- Until convergence (or for a max number of epochs):

$$w_1 := w_1 - \alpha \frac{\partial J(w_1, b)}{\partial w_1} = w_1 - \alpha \frac{1}{N} \sum_{i=0}^N (f_{w_1, b}(x) - y) x$$

$$b := b - \alpha \frac{\partial J(w_1, b)}{\partial b} = b - \alpha \frac{1}{N} \sum_{i=0}^N (f_{w_1, b}(x) - y)$$

Linear Regression: how is J minimized? (GD for 2 params)

$$J(w_1, b) = \frac{1}{2N} \sum_{i=0}^N (f_{w_1, b}(x_i) - y_i)^2 \quad f_{w_1, b} = w_1 x + b$$

Gradient Descent:

- Select a random value for w_1 and b
- Until convergence (or for a max number of epochs):

$$w_1 := w_1 - \alpha \frac{\partial J(w_1, b)}{\partial w_1} = w_1 - \alpha \frac{1}{N} \sum_{i=0}^N (f_{w_1, b}(x) - y) x$$

$$b := b - \alpha \frac{\partial J(w_1, b)}{\partial b} = b - \alpha \frac{1}{N} \sum_{i=0}^N (f_{w_1, b}(x) - y)$$

Why? Just forget Σ and:

$$\begin{aligned} \frac{\partial}{\partial \theta_j} J(\theta) &= \frac{\partial}{\partial \theta_j} \frac{1}{2} (f_{\theta, b}(x) - y)^2 = \\ &= 2 \frac{1}{2} (f_{\theta, b}(x) - y) \frac{\partial}{\partial \theta_j} (f_{\theta, b}(x) - y) = \\ &= (f_{\theta, b}(x) - y) \frac{\partial}{\partial \theta_j} (f_{\theta, b}(x) - y) = \\ &= (f_{\theta, b}(x) - y) \frac{\partial}{\partial \theta_j} (\sum_i \theta_i x_i - y) = \\ &= (f_{\theta, b}(x) - y) x_j \end{aligned}$$

Linear Regression: Gradient Descent Python Example

<https://github.com/tyiannak/ml-python>

See example `1-linear-regression.ipynb`

Linear Regression: how is J minimized? (GD for 2 params)

$$J(w_1, b) = \frac{1}{2N} \sum_{i=0}^N (f_{w_1, b}(x_i) - y_i)^2 \quad f_{w_1, b} = w_1 x + b$$

Gradient Descent:

- Select a random value for w_1 and b
- Until convergence (or for a max number of epochs):

$$w_1 := w_1 - \alpha \frac{\partial J(w_1, b)}{\partial w_1} = w_1 - \alpha \frac{1}{N} \sum_{i=0}^N (f_{w_1, b}(x) - y) x$$

$$b := b - \alpha \frac{\partial J(w_1, b)}{\partial b} = b - \alpha \frac{1}{N} \sum_{i=0}^N (f_{w_1, b}(x) - y)$$

Batch gradient descent: at each step ALL data points are needed

(can be very slow for big data)

Linear Regression: Stochastic Gradient Descent


Stochastic gradient descent

Iteratively uses the derivative of one single example

Goes close to global minimum faster but with a more noisy "path"

Good for very large datasets

(never quite converges)



Batch gradient descent: at each step ALL data points are needed

(can be very slow for big data)

Linear regression: closed-form solution

- **Gradient descent** is an iterative algo that minimizes the cost, by gradually reducing it
- **Particularly for linear regression**, cost minimization can be computed through a closed-form solution
- Compute function's minimum \rightarrow set partial derivatives to zero (critical point): direct solution
- X : $N \times D$ matrix (rows represent training examples)
- y : target values (N -dimensional vector)

$$\mathbf{w} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

Linear regression: why iterative if there's a closed-form solution?

- **Gradient descent** is an iterative algo that minimizes the cost, by gradually reducing it
- Reason: computational complexity: $X^T X$ and (especially) inversion can take a long time
- GD is much more easily parallelizable than matrix manipulation

$$\mathbf{w} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

Regression Evaluation Metrics

- (More details on train/val/test and other ML evaluation issues in later courses)

$$R^2 = 1 - \frac{SS_{Regression}}{SS_{Total}} = 1 - \frac{\sum_i (y_i - \hat{y}_i)^2}{\sum_i (y_i - \bar{y})^2}$$

$$MSE = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$$

$$MAE = \frac{1}{N} \sum_{i=1}^N |y_i - \hat{y}_i|$$